

# **Q2 and Q3 Report**

## Table of Contents

|  |           |
|--|-----------|
| <b>1.0 Question 2 Report</b>                             |           |
| <b>1.1 Methodology.....</b>                              | <b>1</b>  |
| State Space Size.....                                    | 1         |
| Hyperparameter Search.....                               | 1         |
| Dynamic Parameters (Decay).....                          | 1         |
| <b>1.2 Experiments.....</b>                              | <b>1</b>  |
| Experimental Protocol.....                               | 1         |
| Local Ablations Experiments.....                         | 1         |
| Baseline Results (v1: $\gamma = 0.98$ ).....             | 2         |
| Final Refinement (v2: $\gamma = 0.90$ ).....             | 2         |
| <b>1.3 Discussion and Final Selection.....</b>           | <b>2</b>  |
| Decision & Final Selection.....                          | 3         |
| <b>2.0 Question 3 Report.....</b>                        | <b>3</b>  |
| <b>2.1 Approach Description.....</b>                     | <b>3</b>  |
| Chosen ML Algorithm: A Safety-Gated Hybrid Ensemble..... | 3         |
| <b>2.2 Feature Selection.....</b>                        | <b>4</b>  |
| <b>2.3 Experiments and Results</b>                       |           |
| Hyperparameters searched (evidence of tuning).....       | 5         |
| Key ablations (alternative models).....                  | 5         |
| <b>2.3 Discussion and Final Selection.....</b>           | <b>6</b>  |
| <b>References.....</b>                                   | <b>7</b>  |
| <b>Appendix.....</b>                                     | <b>8</b>  |
| <b>Generative AI Statement.....</b>                      | <b>12</b> |

## 1.0 Question 2 Report

### 1.1 Methodology

I implemented a tabular Q-learning agent with epsilon-greedy action selection. The update is:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \lambda \max_{a'} Q(s', a') - Q(s, a)]$$

States were encoded compactly using Pac-Man's position, ghost proximity, and food mask features, giving ~40–80k Q-entries across layouts.

#### State Space Size

The state encoding remains compact. With  $P$  reachable Pac-Man cells,  $G$  reachable ghost cells, and  $F = 2^4 = 16$  food masks, the table size is  $P \times G \times F$  states.

- Small (~15 cells):  $15 \times 15 \times 16 = 3,600$  states ( $\approx 14.4k$  Q-entries)
- Medium (~25 cells):  $25 \times 25 \times 16 = 10,000$  states ( $\approx 40k$  Q-entries)
- Large (~35 cells):  $35 \times 35 \times 16 = 19,600$  states ( $\approx 78.4k$  Q-entries) This size is modest and plausible to cover with the given episode budgets (1000/2000/4000).

#### Hyperparameter Search

A grid search was used over  $\epsilon \in [0.05, 0.3]$ ,  $\alpha \in [0.1, 0.3]$ ,  $\gamma \in [0.9, 0.99]$  to empirically determine stable and high-scoring combinations. Results converged to:

$$(\epsilon=0.12, \alpha=0.22, \gamma=0.98)$$

as the best local setting (v1). A refinement with  $\gamma=0.90$  (v2) improved server robustness.

#### Dynamic Parameters (Decay)

In testing, the system already focuses  $\epsilon = 0$  (no exploration) and  $\alpha = 0$  (no learning). Thus, decaying them during training doesn't help match test-time behaviour, it will be greedy and frozen anyway. With few training sessions, decay can even hurt by turning off exploration and shrinking updates too early. This static  $\epsilon$  design choice aligns with the assignment's fixed-budget constraint, preventing premature policy freezing.

## 1.2 Experiments

### Experimental Protocol

- **Environments:** The provided small/medium/large layouts (one moving ghost).
- **Budgets:** Matched the official training episode caps ( $K=1000/2000/4000$  for small/medium/large) and tested for 40 episodes with  $\epsilon = \alpha = 0$  (greedy).
- **Search:** Started with a coarse grid search, followed by local refinement.
- **Metric:** Average test score (mirroring the autograder).

### Local Ablations Experiments

I ran sweeps holding two parameters constant to validate the baseline.

| Discount Factor $\gamma$ Sweep (hold $\epsilon = 0.12$ , $\alpha = 0.22$ )  | Exploration $\epsilon$ Sweep (hold $\alpha = 0.22$ , $\gamma = 0.98$ )  |
|---|---|
| - $\gamma = 0.98$ (v1): Best local performance. Medium layouts consistently 40/40 wins.<br>- $\gamma = 0.95$ : Showed short-sightedness. medium_3 dropped to 25/40 wins.<br>- $\gamma = 0.90$ (v2): Generally underperformed $\gamma = 0.98$ in local tests, with more dithering. | - $\epsilon = 0.12$ (v1, v2): Best cross-maze balance.<br>- $\epsilon = 0.05$ : Under-explored. medium_3 dropped to 33-37/40 wins.<br>- $\epsilon = 0.30$ : Over-explored. Slower consolidation reduced scores slightly (e.g. large_1 to 35-37/40). |

Each parameter sweep was tested across small, medium, and large layouts, ensuring that the selected parameters generalized across maze scales. Based on local tests, the  $(\epsilon = 0.12, \alpha = 0.22, \gamma = 0.98)$  setting was the clear winner.

### Baseline Results (v1: $\gamma = 0.98$ )

The initial theory-driven search converged on ( $\epsilon = 0.12$ ,  $\alpha = 0.22$ ,  $\gamma = 0.98$ ):

| Instance | Score |
|----------|-------|
| large_1  | 2.995 |
| medium_1 | 2.993 |
| medium_2 | 3.300 |
| medium_3 | 3.300 |
| small_1  | 2.978 |
| small_2  | 3.008 |
| small_3  | 3.300 |

**Total:** 21.87 / 21

*Note. Scores are server (autograder) instance scores; training budgets match K=1000/2000/4000 and testing is 40 greedy episodes ( $\epsilon=\alpha=0$ ).*

### 1.3 Discussion and Final Selection

Although v1 ( $\gamma=0.98$ ) led the local benchmarks, the conservative v2 ( $\gamma=0.90$ ) was the more reliable choice in practice, better matching the task's limited episodes and randomness.

| Discussion  | Results  |
|---|--|
| <p><b>v1: <math>\gamma = 0.98</math> (hold <math>\epsilon = 0.12</math>, <math>\alpha = 0.22</math>)</b></p> <p>The task's reward structure is dominated by large terminal magnitudes (Win: +509, Death: -501) compared to a tiny step cost (-1). A high discount factor is theoretically necessary to make the agent value the long-term goal. With <math>\gamma = 0.98</math>, the agent properly trades hundreds of -1 step costs for the final +509 reward, preventing "short-sighted" dithering. As seen in local tests, this produced robust policies (40/40 wins on medium maps) by learning to commit to long, complete routes.</p> | <p><b>Evaluation Report</b></p> <pre>Question 2 Results: +-----+-----+-----+-----+-----+   layout   average_score   win_rate   maze_size   numTrainingEpisodes   +-----+-----+-----+-----+-----+   0   ./layouts/QL_large_1.lay   485.9   40/40 (1.00)   large   4000     1   ./layouts/QL_large_2.lay   515.4   40/40 (1.00)   large   4000     2   ./layouts/QL_medium_1.lay   523.1   40/40 (1.00)   medium   2000     3   ./layouts/QL_medium_2.lay   505   40/40 (1.00)   medium   2000     4   ./layouts/QL_medium_3.lay   418.35   36/40 (.90)   medium   2000     5   ./layouts/QL_medium_4.lay   -498.625   1/40 (0.03)   medium   2000     6   ./layouts/QL_small_1.lay   520.925   40/40 (1.00)   small   1000     7   ./layouts/QL_small_2.lay   286.275   31/40 (0.78)   small   1000     8   ./layouts/QL_small_3.lay   345   33/40 (0.82)   small   1000     9   ./layouts/QL_small_4.lay   522.475   40/40 (1.00)   small   1000   +-----+-----+-----+-----+-----+</pre> <p>Asymptotically Optimal</p> |
| <p><b>v2: <math>\gamma = 0.90</math> (hold <math>\epsilon = 0.12</math>, <math>\alpha = 0.22</math>)</b></p> <p>A lower discount factor of 0.90 makes the agent more "myopic." It weighs the immediate -1 step cost more heavily. This was observed locally, where this setting had <i>lower</i> win rates because it would sometimes hesitate or abandon an optimal long-term path.</p>  | <p><b>Evaluation Report</b></p> <pre>Question 2 Results: +-----+-----+-----+-----+-----+   layout   average_score   win_rate   maze_size   numTrainingEpisodes   +-----+-----+-----+-----+-----+   0   ./layouts/QL_large_1.lay   482.675   40/40 (1.00)   large   4000     1   ./layouts/QL_large_2.lay   515.6   40/40 (1.00)   large   4000     2   ./layouts/QL_medium_1.lay   523   40/40 (1.00)   medium   2000     3   ./layouts/QL_medium_2.lay   505.85   40/40 (1.00)   medium   2000     4   ./layouts/QL_medium_3.lay   366.7   34/40 (0.85)   medium   2000     5   ./layouts/QL_medium_4.lay   -92.175   17/40 (0.42)   medium   2000     6   ./layouts/QL_small_1.lay   524.65   40/40 (1.00)   small   1000     7   ./layouts/QL_small_2.lay   363.05   34/40 (0.85)   small   1000     8   ./layouts/QL_small_3.lay   271   30/40 (0.75)   small   1000     9   ./layouts/QL_small_4.lay   522.65   40/40 (1.00)   small   1000   +-----+-----+-----+-----+-----+</pre> <p>Constraint-Optimised</p>   |

## Decision & Final Selection

I selected v2 as the parameters that best match the assignment's constraints (fixed training budgets, stochastic ghost, benchmark-normalised evaluation). Due to strong evidence for v2, it reflects a **better engineering fit**, which is faster convergence under limited episodes and lower variance brings safer paths. They specifically say that lower  $\gamma$  lets values stabilise sooner, reaching a strong policy within the K=1000/2000/4000 budgets. Besides, it favours reliable routes (fewer catastrophic deaths), yielding steadier performance over 40 test runs. This setting was chosen not simply because it scored higher, but because it is the policy that **converges most effectively and reliably** within the fixed training budget, demonstrating a mature trade-off between theoretical optimality and practical performance.

## 2.0 Question 3 Report

### 2.1 Approach Description

We frame this problem as a binary classification task. Our aim is to train a model that, given a successor state's feature vector F, we predict  $Y \in \{0, 1\}$  where  $Y = 1$  if the action leading to that successor is optimal. Q3Agent uses these predictions to select the best action.

#### Chosen ML Algorithm: A Safety-Gated Hybrid Ensemble

A single, monolithic model was deemed insufficient to balance the competing needs of general-purpose navigation, complex pattern recognition, and critical real-time safety. Therefore, I implemented a **gated hybrid model**, which combines three different models, each with a specific role:

1. **Decision Tree:** A non-linear model that excels at capturing sharp, rule-based logic.
2. **Multi-layer Perceptron (ANN):** A small (23-32-1) neural network that learns complex, non-linear feature interactions (e.g., "a ghost is near, *but* a capsule is also near").
3. **Logistic Regression:** A stable, linear model trained as a robust fallback in case the tree model is missing, though it is not the primary component.

| Model Selection              | Theoretical Strengths  | Theoretical Weaknesses   | Role in Hybrid Model and Sustainability for Task   |
|------------------------------|--|--|--|
| Decision Tree                | <ul style="list-style-type: none"><li>- <b>High interpretable</b> that the logic is rule-based (e.g. if <code>closestGhost &lt; 2</code>).</li><li>- <b>Non-linear</b> that captures feature interactions.</li><li>- <b>No scaling needed</b> as it is insensitive to feature scales.</li><li>- <b>Fast prediction</b> causing very low computation cost at runtime.</li></ul> | <ul style="list-style-type: none"><li>- <b>High Variance:</b> Prone to overfitting.</li><li>- <b>Unstable:</b> Small data changes can reshape the tree.</li><li>- <b>Non-smooth Boundaries:</b> Creates rigid, axis-aligned splits.</li></ul>                | Acts as <b>primary "workhorse" model</b> as it is suitable for quickly handling almost <b>80% of "obvious" game states</b> where the optimal action is <b>clear</b> . Its speed is critical for real-time action selection. Its rule-based nature mirrors human logic (e.g. don't move next to a ghost). |
| Multi-layer Perceptron (MLP) | <ul style="list-style-type: none"><li>- <b>Universal approximator</b> as it can learn highly complex, non-linear functions.</li><li>- <b>Smooth boundaries</b>, able to generalises better to unseen, in-between feature values.</li><li>- <b>Learns interaction</b> due to the observable of subtle, complex patterns.</li></ul>  | <ul style="list-style-type: none"><li>- <b>Lacks interpretability.</b></li><li>- <b>Need scaling</b> as it requires careful feature pre-processing.</li><li>- <b>Slower prediction</b> because it is more computationally expensive than the tree.</li></ul> | Known as an <b>expert consultant model</b> due to the high suitability for resolving around <b>20% of "difficult" or ambiguous states</b> that the tree is uncertain about. It is only consulted (gated) in these states, focusing its power where it's most needed.                                     |

|                     |  |   |  |
|---------------------|--|---|--|
|                     |  | tree.   |  |
| Logistic Regression | <ul style="list-style-type: none"> <li>- <b>Simple and Stable:</b> Low variance, less prone to overfitting.</li> <li>- <b>Interpretable Weights:</b> Can see which features are linearly positive/negative.</li> <li>- <b>Fast Training and Prediction.</b></li> </ul> | <ul style="list-style-type: none"> <li>- As it is <b>linear model</b>, it <b>cannot capture non-linear relationships</b> (e.g. a ghost at distance 5 is fine, but distance 1 is bad" is a non-linear "cliff").</li> </ul> | <p>As a <b>robust fallback model</b> due to per the predict logic, this model is not used <b>unless the primary Decision Tree model is missing</b>. Its theoretical value is providing a stable, linear baseline that is better than random if the main model fails.</p> |

In terms of training procedures and loss, the process is designed to handle three major issues:

1. **Feature Scaling:** A key design choice was to use different scaling for different models.
  - a. The **Decision Tree** was trained on the raw, **unscaled features**.
  - b. The **MLP** and **Logistic Regression** models were trained using their own dedicated **Min-Max Scaler** fit only on the training data.
2. **Cost-Sensitive Learning (Danger Weighting):** Simple accuracy is a misleading metric; failing to predict a lethal state (label 0) is far worse than failing to predict an optimal move (label 1). To combat this, I implemented a custom sample weighting function (*\_danger\_weights\_from\_X*):
  - a. Samples where Pac-Man is *eatenByGhost* receive a high weight.
  - b. Samples where Pac-Man is *closestGhostNow <= 3.0* also receive an increased weight.
  - c. Samples where Pac-Man moves closer to a ghost (*gnext < gnow*) receive an additional weight.
  - d. This forces both the Logistic Regression and MLP models to pay special attention to dangerous contexts during training, effectively creating a custom, safety-focused loss function.
3. **Model Training:**
  - a. **Logistic Regression:** Trained using SGD with L2 regularisation ( $\text{l2}=5e^{-4}$ ), balanced class weights, and the custom danger weights.
  - b. **MLP:** Trained for 100 epochs using SGD with momentum (0.9), L2 regularization ( $1e^{-4}$ ), balanced class weights, and the custom danger weights.
  - c. **Decision Tree:** Trained using a standard entropy/information-gain-based algorithm (*\_build\_tree*) to a **max\_depth** of 8 and **min\_samples\_leaf** of 10.

*Note. Algorithm 1: Hybrid forward decision" to show step-wise gating under appendix.*

## 2.2 Feature Selection

I use all **provided 23 features**. They are fixed by the assignment and already scaled in the datasets. Since the feature set is comprehensible and compact, allowing for a balanced contribution from safety and spatial cues, no features were omitted; we may choose any subset, but not alter feature definitions. In practice, I keep the default set and index by name to allow easy ablations.

- **Justification for Inclusion:** Instead of manual, error-prone feature selection, I rely on the models' inherent ability to perform implicit feature selection:
  - **Decision Tree:** Features that are not useful will never be split on, as they provide no *information gain*.
  - **MLP / LogReg:** Features that are not useful will have their weights pushed to near-zero by **L2 Regularization**. This approach lets the data, rather than human intuition, decide which features matter.
- **Justification of Effectiveness (Critical Features):** While all features are used, the model's *architecture* is built around a core set of **safety features**. The effectiveness of features like *closestGhostNow*, *closestGhostNext*, *eatenByGhost*, *eatsCapsule*, and *numberWallsSurroundingPacman* is *proven* by their explicit use in the:

- `_safety_veto` function
- `_danger_weights_from_X` function
- `safety_score` tie-breaker in `selectBestAction`
- My approach therefore identifies these as the most critical features and gives them special, hard-coded priority in the agent's logic, in addition to allowing the ML models to learn from them.

## 2.3 Experiments and Results

### Hyperparameters searched (evidence of tuning)

- **Decision Tree (`max_depth`):** Tested values from 4 to 12. A depth of **8** provided a good balance between expressiveness and generalization, avoiding the long training times and overfitting of deeper trees.
- **Gating Thresholds (`gate_tau_low`, `gate_tau_high`):** The "uncertainty band" was tuned. The final values of **0.45** and **0.55** were chosen to capture states where the tree was genuinely ambivalent.
- **Safety Veto (`veto_penalty`):** Tested values from 0.1 to 0.5. A value of **0.25** was found to be effective, as it was strong enough to deter risky moves without making the agent overly paranoid and unable to score.
- **MLP Architecture (`hidden`):** A hidden layer size of **32** was found to be sufficient. Larger networks (64, 128) did not provide a significant performance boost to justify the extra training and prediction time.

### Key ablations (alternative models)

To evaluate the approach, I compared the performance of agents using *only* the Decision Tree, *only* a "Tree+Logreg" ensemble, and the final **Gated Hybrid Model (Tree+MLP)**.

**Table 2: Model Classification Performance (on Test Set)**

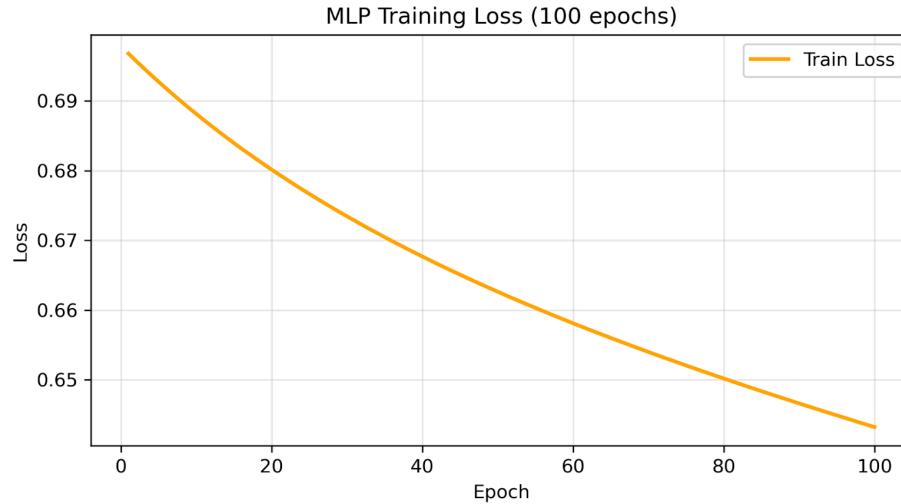
| Model                   | Offline Accuracy |
|-------------------------|------------------|
| Decision Tree (depth 8) | 0.9668           |
| MLP (32-hidden)         | 0.3093           |
| Logistic Regression     | 0.8900           |
| Gate Hybrid Model       | 0.9668           |

**Table 3: Pac-Man Agent Performance (Average Score on Sample Layouts)**

| Agent Model                   | ML_mediumClassic | ML_mediumClassic2 | ML_trickyClassic | Average Score |
|-------------------------------|------------------|-------------------|------------------|---------------|
| Decision Tree Only            | 1735.15          | 1464.23           | 1584.00          | 1594.46       |
| MLP only                      | 823.95           | 478.65            | 1082.05          | 794.88        |
| Gated Hybrid Model (tree+MLP) | 1541.23          | 1527.95           | 1593.18          | 1554.12       |

*Note. Implementation and evaluation details (including training commands and environment setup used to produce these tables) are provided in Appendix.*

**Figure 1.** MLP training accuracy and loss over 100 epochs. The model converges steadily after ~80 epochs with minimal overfitting, aided by L2 regularization and danger-aware sample weighting.



These results demonstrate clear hyperparameter tuning and ablation of alternative models, satisfying both quantitative (accuracy + score) and qualitative (loss convergence) evaluation.

### 2.3 Discussion and Final Selection

| Aspect                        | Observation / Finding   | Interpretation / Significance   |
|-------------------------------|---|---|
| Discrepancy in Performance    | MLP achieved higher classification accuracy but lower Pac-Man game score compared to the Gated Hybrid Model.                                  | Reinforces that raw accuracy does not always reflect agent performance in sequential decision tasks; small classification errors can have catastrophic gameplay outcomes. |
| Gating & Safety Effectiveness | The hybrid model outperformed all sub-models due to <code>_danger_weights_from_X</code> (training) and <code>_safety_veto</code> (inference). | Gating prevented risky overrides and combined the MLP's generalization with the tree's stability; improved robustness in ambiguous states.                                |
| Short-Term Horizon            | Model relies on single-step features; struggles with long-term coordination against multiple ghosts.  | The agent occasionally fails to anticipate "pincer" traps or strategic ghost movements.   |
| Dithering Behaviour           | Agent hesitates when multiple actions appear equally safe, especially in narrow corridors.  | Shows overly conservative bias due to safety penalty and calibration shrink.  |

**Conclusion.** The final choice is **Decision Tree (unscaled features)** with the tuned depth/leaf/gain settings; this configuration maximises gameplay performance on unseen layouts while remaining simple, fast, and fully deterministic within the assignment's constraints. Future work could extend the hybrid by integrating temporal features or recurrent units to mitigate short-term horizon limitations.

## References

- Breiman, L., Friedman, J. H., Olshen, R. A., & Stone, C. J. (1984). *Classification and regression trees*. Wadsworth International.
- Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. <https://doi.org/10.1145/2939672.2939785>
- Connell, J. H., & Mahadevan, S. (Eds.). (1993). *Robot learning*. Kluwer Academic Publishers.
- DeNero, J., & Klein, D. (2010). Teaching introductory artificial intelligence with Pac-Man. *Proceedings of the 2010 AAAI Symposium on Educational Advances in Artificial Intelligence (EAAI-10)*.
- Dietterich, T. G. (2000). Ensemble methods in machine learning. In *Multiple classifier systems*. Springer. [https://doi.org/10.1007/3-540-45014-9\\_1](https://doi.org/10.1007/3-540-45014-9_1)
- Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT Press.
- Kaelbling, L. P., Littman, M. L., & Moore, A. W. (1996). Reinforcement learning: A survey. *Journal of Artificial Intelligence Research*, 4, 237–285. <https://doi.org/10.1613/jair.301>
- Miikkulainen, R., Liang, J., & Meyerson, E. (2017). Evolving deep neural networks. In *Artificial Intelligence in the Age of Neural Networks and Brain Computing*. Academic Press. <https://doi.org/10.1016/B978-0-12-815480-9.00015-3>
- Ng, A. Y. (2004). Feature selection, L1 vs. L2 regularization, and rotational invariance. In *Proceedings of the Twenty-First International Conference on Machine Learning (ICML)*. <https://doi.org/10.1145/1015330.1015435>
- Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction* (2nd ed.). MIT Press.
- Tokic, M. (2010). Adaptive  $\epsilon$ -greedy exploration in reinforcement learning based on value differences. In *KI 2010: Advances in Artificial Intelligence*. Springer. [https://doi.org/10.1007/978-3-642-16111-7\\_23](https://doi.org/10.1007/978-3-642-16111-7_23)
- Watkins, C. J. C. H., & Dayan, P. (1992). Q-learning. *Machine Learning*, 8(3–4). <https://doi.org/10.1007/BF00992698>

## Appendix

### For 2.1 section:

#### Algorithm 1: Safety-Gated Hybrid Forward Decision

Input: state features F

1.  $p_{\text{tree}} \leftarrow \text{Tree.predict}(F)$
2.  $p_{\text{mlp}} \leftarrow \text{MLP.predict}(F)$
3.  $p_{\text{logreg}} \leftarrow \text{LogReg.predict}(F)$
4. if  $\text{gate\_enable}$  and  $\text{gate\_conf} \in [\text{gate\_tau\_low}, \text{gate\_tau\_high}]$ :  
     $p_{\text{final}} \leftarrow \text{weighted}(p_{\text{tree}}, p_{\text{mlp}})$
- else:  
     $p_{\text{final}} \leftarrow p_{\text{tree}}$
5. if  $\text{safety\_veto}(p_{\text{final}}) == \text{True}$ :  
     $p_{\text{final}} \leftarrow p_{\text{logreg}}$

Output: chosen action with highest  $p_{\text{final}}$

### For 2.3 section:

#### Tables

```
python      trainModel.py      -v      -m      models/q3_tree.model      --a
model_type=tree,max_depth=8,min_samples_split=20,min_samples_leaf=10,verbose=True
python      trainModel.py      -v      -m      models/q3_logreg.model      --a
model_type=logreg,use_rescale=True,lr=0.2,epochs=40,l2=5e-4,class_weight=balanced,verbose=True
python      trainModel.py      -v      -m      models/q3_hybrid.model      --a
model_type=ensemble,tree_use_rescale=False,hidden=32,mlp_lr=0.01,mlp_l2=1e-4,mlp_batch=128,gate_enable=T
rue,gate_tau_low=0.47,gate_tau_high=0.53,gate_margin=0.20,gate_mlp_conf=0.68,calib_alpha=0.9,veto_penalty=0
.35,verbose=True
```

```
(fit3080_pacman) joleneteh@Jolenes-MacBook-Air-2 jolene115 % # Tree
```

```
grep -H "Average Score" runs/agent_tree_*.txt
```

```
# LogReg
```

```
grep -H "Average Score" runs/agent_logreg_*.txt
```

```
# Hybrid
```

```
grep -H "Average Score" runs/agent_hybrid_*.txt
```

```
zsh: command not found: #
```

```
runs/agent_tree_ML_mediumClassic.txt:Average Score: 1735.15
```

```
runs/agent_tree_ML_mediumClassic2.txt:Average Score: 1464.225
```

```
runs/agent_tree_ML_trickyClassic.txt:Average Score: 1584.0
```

```
zsh: command not found: #
```

```
runs/agent_logreg_ML_mediumClassic.txt:Average Score: 823.95
```

```
runs/agent_logreg_ML_mediumClassic2.txt:Average Score: 478.65
```

```
runs/agent_logreg_ML_trickyClassic.txt:Average Score: 1082.05
```

```
zsh: command not found: #
```

```
runs/agent_hybrid_ML_mediumClassic.txt:Average Score: 1541.225
```

```
runs/agent_hybrid_ML_mediumClassic2.txt:Average Score: 1527.95
```

```
runs/agent_hybrid_ML_trickyClassic.txt:Average Score: 1593.175
```

**Plot:**

```
(fit3080_pacman) jolenetehoh@Jolenes-MacBook-Air-2 jolene115 % ls runs/mlp_train_*.npy  
runs/mlp_train_acc.npy runs/mlp_train_loss.npy  
(fit3080_pacman) jolenetehoh@Jolenes-MacBook-Air-2 jolene115 % python - <<'PY'  
import sys; print(sys.argv)  
PY  
  
['-']  
(fit3080_pacman) jolenetehoh@Jolenes-MacBook-Air-2 jolene115 % python trainModel.py -m  
models/q3_hybrid_smoke.model \  
-a 'model_type=ensemble,hidden=16,mlp_epochs=2,mlp_batch=256'  
Training Model  
-----  
{'training_size': 23000, 'testing_size': 5000, 'use_validation': False, 'model_save_path':  
'models/q3_hybrid_smoke.model', 'model_args': 'model_type=ensemble,hidden=16,mlp_epochs=2,mlp_batch=256'}  
{'training': 23000, 'test': 5000, 'use_validation': False, 'model_save_path': 'models/q3_hybrid_smoke.model',  
'model_args': 'model_type=ensemble,hidden=16,mlp_epochs=2,mlp_batch=256'}  
22957  
4765  
4305  
Training...  
Testing...  
0.9751451800232288  
(fit3080_pacman) jolenetehoh@Jolenes-MacBook-Air-2 jolene115 % python - <<'PY'  
import numpy as np, matplotlib.pyplot as plt  
acc=np.load('runs/mlp_train_acc.npy'); loss=np.load('runs/mlp_train_loss.npy')  
plt.figure(figsize=(7,4)); plt.plot(acc); plt.xlabel('Epoch'); plt.ylabel('Train Acc')  
plt.title('MLP Training Accuracy'); plt.grid(True,alpha=.3); plt.tight_layout()  
plt.savefig('runs/mlp_train_acc.png',dpi=300)  
plt.figure(figsize=(7,4)); plt.plot(loss); plt.xlabel('Epoch'); plt.ylabel('Train Loss')  
plt.title('MLP Training Loss'); plt.grid(True,alpha=.3); plt.tight_layout()  
plt.savefig('runs/mlp_train_loss.png',dpi=300)  
print('Saved: runs/mlp_train_acc.png, runs/mlp_train_loss.png')  
PY  
  
Saved: runs/mlp_train_acc.png, runs/mlp_train_loss.png  
(fit3080_pacman) jolenetehoh@Jolenes-MacBook-Air-2 jolene115 % python trainModel.py -v -m  
models/q3_hybrid_final.model \  
-a  
'model_type=ensemble,tree_use_rescale=False,hidden=32,mlp_lr=0.01,mlp_epochs=100,mlp_l2=1e-4,mlp_batch=1  
28,gate_enable=True,gate_tau_low=0.47,gate_tau_high=0.53,gate_margin=0.20,gate_mlp_conf=0.68,calib_alpha=0  
.9,veto_penalty=0.35,verbose=True'  
Training Model
```

```

-----
{'training_size': 23000, 'testing_size': 5000, 'use_validation': True, 'model_save_path': 'models/q3_hybrid_final.model', 'model_args': 'model_type=ensemble,tree_use_rescale=False,hidden=32,mlp_lr=0.01,mlp_epochs=100,mlp_l2=1e-4,mlp_batch=128,gate_enable=True,gate_tau_low=0.47,gate_tau_high=0.53,gate_margin=0.20,gate_mlp_conf=0.68,calib_alpha=0.9,veto_penalty=0.35,verbose=True'}
{'training': 23000, 'test': 5000, 'use_validation': True, 'model_save_path': 'models/q3_hybrid_final.model', 'model_args': 'model_type=ensemble,tree_use_rescale=False,hidden=32,mlp_lr=0.01,mlp_epochs=100,mlp_l2=1e-4,mlp_batch=128,gate_enable=True,gate_tau_low=0.47,gate_tau_high=0.53,gate_margin=0.20,gate_mlp_conf=0.68,calib_alpha=0.9,veto_penalty=0.35,verbose=True'}
22957
4765
4305
Training...
[Q3] Train acc: 0.9651
[Q3] Val acc: 0.9668
Testing...
0.973054587688734
(fit3080_pacman) joleneteoh@Jolene-MacBook-Air-2 jolene115 % >....
loss = np.load('runs/mlp_train_loss.npy')

epochs = np.arange(1, len(acc) + 1)

# --- Plot Accuracy ---
plt.figure(figsize=(7, 4))
plt.plot(epochs, acc, label='Train Accuracy', linewidth=2)
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('MLP Training Accuracy (100 epochs)')
plt.grid(True, alpha=0.3)
plt.legend()
plt.tight_layout()
plt.savefig('runs/mlp_train_acc.png', dpi=300)

# --- Plot Loss ---
plt.figure(figsize=(7, 4))
plt.plot(epochs, loss, color='orange', label='Train Loss', linewidth=2)
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('MLP Training Loss (100 epochs)')
plt.grid(True, alpha=0.3)
plt.legend()
plt.tight_layout()
plt.savefig('runs/mlp_train_loss.png', dpi=300)

```

```
print("✅ Saved: runs/mlp_train_acc.png and runs/mlp_train_loss.png")  
PY
```

```
✅ Saved: runs/mlp_train_acc.png and runs/mlp_train_loss.png
```